

# L1: NLP tasks with a simple interface

Load your HF API key and relevant Python libraries.

```
In [*]: import os
import io
from IPython.display import Image, display, HTML
from PIL import Image
import base64
from dotenv import load_dotenv, find_dotenv
_ = load_dotenv(find_dotenv()) # read local .env file
hf_api_key = os.environ['HF_API_KEY']
```

```
In [2]: # Helper function
import requests, json

#Summarization endpoint
def get_completion(inputs, parameters=None, ENDPOINT_URL=os.environ['HF_API_SUMMARY_BA
    headers = {
        "Authorization": f"Bearer {hf_api_key}",
        "Content-Type": "application/json"
    }
    data = { "inputs": inputs }
    if parameters is not None:
        data.update({"parameters": parameters})
    response = requests.request("POST",
                                ENDPOINT_URL, headers=headers,
                                data=json.dumps(data)
                                )
    return json.loads(response.content.decode("utf-8"))
```

## Building a text summarization app

Here we are using an [Inference Endpoint \(https://huggingface.co/inference-endpoints\)](https://huggingface.co/inference-endpoints) for the shleifer/distilbart-cnn-12-6 , a 306M parameter distilled model from facebook/bart-large-cnn .

## How about running it locally?

The code would look very similar if you were running it locally instead of from an API. The same is true for all the models in the rest of the course, make sure to check the [Pipelines \(https://huggingface.co/docs/transformers/main\\_classes/pipelines\)](https://huggingface.co/docs/transformers/main_classes/pipelines) documentation page

```
from transformers import pipeline
```

```
get_completion = pipeline("summarization", model="shleifer/distilbart-cnn-12-6")
```

```
def summarize(input):
```

```
    output = get_completion(input)
```

```
In [3]: text = ('''The tower is 324 metres (1,063 ft) tall, about the same height
                as an 81-storey building, and the tallest structure in Paris.
                Its base is square, measuring 125 metres (410 ft) on each side.
                During its construction, the Eiffel Tower surpassed the Washington
                Monument to become the tallest man-made structure in the world,
                a title it held for 41 years until the Chrysler Building
                in New York City was finished in 1930. It was the first structure
                to reach a height of 300 metres. Due to the addition of a broadcasting
                aerial at the top of the tower in 1957, it is now taller than the
                Chrysler Building by 5.2 metres (17 ft). Excluding transmitters, the
                Eiffel Tower is the second tallest free-standing structure in France
                after the Millau Viaduct.''' )

                get_completion(text)
```

```
[{'summary_text': ' The tower is 324 metres (1,063 ft) tall, about the same height as
an 81-storey building . It is the tallest structure in Paris and the second tallest fr
ee-standing structure in France after the Millau Viaduct . It was the first structure
in the world to reach a height of 300 metres .'}]
```

## Getting started with Gradio `gr.Interface`

### How about running it locally?

The code would look very similar if you were running it locally. Simply remove all the paramters in the launch method

```
demo.launch()
```

---

```
In [4]: import gradio as gr
def summarize(input):
    output = get_completion(input)
    return output[0]['summary_text']

gr.close_all()
demo = gr.Interface(fn=summarize, inputs="text", outputs="text")
demo.launch(share=True, server_port=int(os.environ['PORT1']))
```

---

Running on local URL: <https://0.0.0.0:39648> (<https://0.0.0.0:39648>)

Could not create share link. Missing file: /usr/local/lib/python3.9/site-packages/gradio/frpc\_linux\_amd64\_v0.2.

Please check your internet connection. This can happen if your antivirus software blocks the download of this file. You can install manually by following these steps:

1. Download this file: [https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64) ([https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64))
2. Rename the downloaded file to: frpc\_linux\_amd64\_v0.2
3. Move the file to this location: /usr/local/lib/python3.9/site-packages/gradio

## 502 Bad Gateway

---

openresty

You can add `demo.launch(share=True)` to create a public link to share with your team or friends.

```
In [5]: import gradio as gr

def summarize(input):
    output = get_completion(input)
    return output[0]['summary_text']

gr.close_all()
demo = gr.Interface(fn=summarize,
                    inputs=[gr.Textbox(label="Text to summarize", lines=6)],
                    outputs=[gr.Textbox(label="Result", lines=3)],
                    title="Text summarization with distilbart-cnn",
                    description="Summarize any text using the `shleifer/distilbart-cn
                    )
demo.launch(share=True, server_port=int(os.environ['PORT2']))
```

Closing server running on port: 39648

Running on local URL: <https://0.0.0.0:47506> (<https://0.0.0.0:47506>)

Could not create share link. Missing file: /usr/local/lib/python3.9/site-packages/gradio/frpc\_linux\_amd64\_v0.2.

Please check your internet connection. This can happen if your antivirus software blocks the download of this file. You can install manually by following these steps:

1. Download this file: [https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64) ([https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64))
2. Rename the downloaded file to: frpc\_linux\_amd64\_v0.2
3. Move the file to this location: /usr/local/lib/python3.9/site-packages/gradio

# 502 Bad Gateway

---

openresty

## Building a Named Entity Recognition app

We are using this [Inference Endpoint \(https://huggingface.co/inference-endpoints\)](https://huggingface.co/inference-endpoints) for `dslim/bert-base-NER`, a 108M parameter fine-tuned BART model on the NER task.

### How about running it locally?

```
from transformers import pipeline

get_completion = pipeline("ner", model="dslim/bert-base-NER")

def ner(input):
    output = get_completion(input)
    return {"text": input, "entities": output}
```

```
In [6]: API_URL = os.environ['HF_API_NER_BASE'] #NER endpoint
text = "My name is Andrew, I'm building DeepLearningAI and I live in California"
get_completion(text, parameters=None, ENDPOINT_URL= API_URL)
```

```
[{'entity': 'B-PER',
  'score': 0.9990625,
  'index': 4,
  'word': 'Andrew',
  'start': 11,
  'end': 17},
 {'entity': 'B-ORG',
  'score': 0.9927857,
  'index': 10,
  'word': 'Deep',
  'start': 32,
  'end': 36},
 {'entity': 'I-ORG',
  'score': 0.99677867,
  'index': 11,
  'word': '##L',
  'start': 36,
  'end': 37},
 {'entity': 'I-ORG',
  'score': 0.9954496,
  'index': 12,
  'word': '##ear',
  'start': 37,
  'end': 40},
 {'entity': 'I-ORG',
  'score': 0.9959293,
  'index': 13,
  'word': '##ning',
  'start': 40,
  'end': 44},
 {'entity': 'I-ORG',
  'score': 0.8917463,
  'index': 14,
  'word': '##A',
  'start': 44,
  'end': 45},
 {'entity': 'I-ORG',
  'score': 0.50361204,
  'index': 15,
  'word': '##I',
  'start': 45,
  'end': 46},
 {'entity': 'B-LOC',
  'score': 0.99969244,
  'index': 20,
  'word': 'California',
  'start': 61,
  'end': 71}]
```

```
In [7]: def ner(input):
        output = get_completion(input, parameters=None, ENDPOINT_URL=API_URL)
        return {"text": input, "entities": output}

        gr.close_all()
        demo = gr.Interface(fn=ner,
                            inputs=[gr.Textbox(label="Text to find entities", lines=2)],
                            outputs=[gr.HighlightedText(label="Text with entities")],
                            title="NER with dslim/bert-base-NER",
                            description="Find entities using the `dslim/bert-base-NER` model",
                            allow_flagging="never",
                            #Here we introduce a new tag, examples, easy to use examples for .
                            examples=["My name is Andrew and I live in California", "My name
demo.launch(share=True, server_port=int(os.environ['PORT3']))
```

Closing server running on port: 39648

Closing server running on port: 47506

Running on local URL: <https://0.0.0.0:41617> (<https://0.0.0.0:41617>)

Could not create share link. Missing file: /usr/local/lib/python3.9/site-packages/grad  
io/frpc\_linux\_amd64\_v0.2.

Please check your internet connection. This can happen if your antivirus software bloc  
ks the download of this file. You can install manually by following these steps:

1. Download this file: [https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64) ([https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc\\_linux\\_amd64](https://cdn-media.huggingface.co/frpc-gradio-0.2/frpc_linux_amd64))
2. Rename the downloaded file to: frpc\_linux\_amd64\_v0.2
3. Move the file to this location: /usr/local/lib/python3.9/site-packages/gradio



# 502 Bad Gateway

---

openresty

## Adding a helper function to merge tokens

```
In [*]: def merge_tokens(tokens):
    merged_tokens = []
    for token in tokens:
        if merged_tokens and token['entity'].startswith('I-') and merged_tokens[-1]['entity'] == token['entity']:
            # If current token continues the entity of the last one, merge them
            last_token = merged_tokens[-1]
            last_token['word'] += token['word'].replace('##', '')
            last_token['end'] = token['end']
            last_token['score'] = (last_token['score'] + token['score']) / 2
        else:
            # Otherwise, add the token to the list
            merged_tokens.append(token)

    return merged_tokens

def ner(input):
    output = get_completion(input, parameters=None, ENDPOINT_URL=API_URL)
    merged_tokens = merge_tokens(output)
    return {"text": input, "entities": merged_tokens}

gr.close_all()
demo = gr.Interface(fn=ner,
                    inputs=[gr.Textbox(label="Text to find entities", lines=2)],
                    outputs=[gr.HighlightedText(label="Text with entities")],
                    title="NER with dsllm/bert-base-NER",
                    description="Find entities using the `dsllm/bert-base-NER` model",
                    allow_flagging="never",
                    examples=["My name is Andrew, I'm building DeeplearningAI and I l"])

demo.launch(share=True, server_port=int(os.environ['PORT4']))
```

Closing server running on port: 39648

Closing server running on port: 47506

Closing server running on port: 41617

Running on local URL: <https://0.0.0.0:57171> (<https://0.0.0.0:57171>)

```
In [*]: gr.close_all()
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

