

L2: Image captioning app

Load your HF API key and relevant Python libraries

```
In [ ]: import os
import io
import IPython.display
from PIL import Image
import base64
from dotenv import load_dotenv, find_dotenv
_ = load_dotenv(find_dotenv()) # read local .env file
hf_api_key = os.environ['HF_API_KEY']
```

```
In [ ]: # Helper functions
import requests, json

#Image-to-text endpoint
def get_completion(inputs, parameters=None, ENDPOINT_URL=os.environ['HF_API_ITT_BASE'],
    headers = {
        "Authorization": f"Bearer {hf_api_key}",
        "Content-Type": "application/json"
    },
    data = { "inputs": inputs },
    if parameters is not None:
        data.update({"parameters": parameters})
    response = requests.request("POST",
                                ENDPOINT_URL,
                                headers=headers,
                                data=json.dumps(data))
    return json.loads(response.content.decode("utf-8"))
```

Building an image captioning app

Here we'll be using an [Inference Endpoint \(https://huggingface.co/inference-endpoints\)](https://huggingface.co/inference-endpoints) for Salesforce/blip-image-captioning-base a 14M parameter captioning model.

The code would look very similar if you were running it locally instead of from an API. You can check the [Pipelines \(https://huggingface.co/docs/transformers/main_classes/pipelines\)](https://huggingface.co/docs/transformers/main_classes/pipelines) documentation page.

```
from transformers import pipeline
```

```
get_completion(pipeline("image-to-text", model="Salesforce/blip-image-captioning-base"))
```

The free images are available on: <https://free-images.com/> (<https://free-images.com/>)

```
In [ ]: image_url = "https://free-images.com/sm/9596/dog_animal_greyhound_983023.jpg"
display(IPython.display.Image(url=image_url))
get_completion(image_url)
```

Captioning with gr.Interface()

```
In [ ]: import gradio as gr

def image_to_base64_str(pil_image):
    byte_arr = io.BytesIO()
    pil_image.save(byte_arr, format='PNG')
    byte_arr = byte_arr.getvalue()
    return str(base64.b64encode(byte_arr).decode('utf-8'))

def captioner(image):
    base64_image = image_to_base64_str(image)
    result = get_completion(base64_image)
    return result[0]['generated_text']

gr.close_all()
demo = gr.Interface(fn=captioner,
                    inputs=[gr.Image(label="Upload image", type="pil")],
                    outputs=[gr.Textbox(label="Caption")],
                    title="Image Captioning with BLIP",
                    description="Caption any image using the BLIP model",
                    allow_flagging="never",
                    examples=["christmas_dog.jpeg", "bird_flight.jpeg", "cow.jpeg"])

demo.launch(share=True, server_port=int(os.environ['PORT1']))
```

```
In [ ]: gr.close_all()
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: